

ADVANCED DATA SCIENCE


PART 2

ALEXANDRE GRAMFORT
THOMAS MOREAU

ENSEMBLE OF EXPERTS: BOOSTING

- Each model is an expert on the errors of its predecessor
- Iteratively re-weights training examples based on errors

- ERM with weights:
$$\arg \min_{f \in \mathcal{F}} \frac{1}{n} \sum_i w_i \ell(f(x_i), y_i)$$



ADABOOST [Y. FREUND & R. SCHAPIRE, 1995]

ADABOOST($D_n = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$, BASE(\cdot, \cdot), T)

For binary classification

```
1   $\mathbf{w}^{(1)} \leftarrow (1/n, \dots, 1/n)$   $\triangleright$  initial weights
2  for  $t \leftarrow 1$  to  $T$ 
3       $h^{(t)} \leftarrow \text{BASE}(D_n, \mathbf{w}^{(t)})$   $\triangleright$  calling the base learner
4       $\gamma^{(t)} \leftarrow \sum_{i=1}^n w_i^{(t)} h^{(t)}(\mathbf{x}_i) y_i$   $\triangleright$  edge = 1 - 2 × error
5       $\alpha^{(t)} \leftarrow \frac{1}{2} \ln \left( \frac{1 + \gamma^{(t)}}{1 - \gamma^{(t)}} \right)$   $\triangleright$  coefficient of  $h^{(t)}$ 
6      for  $i \leftarrow 1$  to  $n$   $\triangleright$  re-weighting the points
7          if  $h^{(t)}(\mathbf{x}_i) \neq y_i$  then
8               $w_i^{(t+1)} \leftarrow w_i^{(t)} \frac{1}{1 - \gamma^{(t)}}$ 
9          else
10              $w_i^{(t+1)} \leftarrow w_i^{(t)} \frac{1}{1 + \gamma^{(t)}}$ 
11 return  $f^{(T)}(\cdot) = \sum_{t=1}^T \alpha^{(t)} h^{(t)}(\cdot)$ 
```

Remark: Freund & Schapire won the Gödel prize 2003

GRADIENT BOOSTING

- Gradient Boosting generalizes adaboost to any arbitrary loss
- a.k.a. GB(R)T, Gradient boosting (regression) trees
- It was originally proposed by [J. Friedman, 1999]
- Variants of the original GBT algorithm are now state-of-the-art models.
- Numerous successes in Kaggle competitions
- State-of-the-art implementations:
 - XGBoost [Chen & Guestrin, Arxiv 2016] (w. Apple, NVidia)
 - LightGBM [Ke et al., Proc. NIPS 2017] (by Microsoft)
 - CatBoost [Prokhorenkova et al. Arxiv 2017] (by Yandex)
 - `sklearn.ensemble.HistGradientBoostingClassifier` (v0.21)

TREE BOOSTING IN A NUTSHELL

$$\hat{y}_i = \sum_{t=1}^T f_t(x_i), f_t \in \mathcal{F} \quad (\text{additive ensemble model})$$

where $\mathcal{F} = \{f(x) = z_{q(x)}^f\}$ (set of trees with K leafs)

with $q : \mathbb{R}^m \rightarrow \{1, \dots, K\}$ (tree partitioning structure)

and $z^f \in \mathbb{R}^K$ (leaf values)

Each f_t has a different tree structure

Remark: Continuous values even for classification

based on [Chen & Guestrin, Arxiv 2016]

TREE BOOSTING IN A NUTSHELL

Objective function: (smooth convex loss function)

$$\mathcal{L}(f_1, \dots, f_T) = \sum_i \ell(y_i, \hat{y}_i) + \sum_t \Omega(f_t)$$

Where: $\Omega(f) = \gamma K + \frac{\lambda}{2} \|z^f\|^2$ (penalize complex trees)

Remark: Original GBT algo. by Friedman had no regularization

Example with MSE / L2 loss:

$$\mathcal{L}(f_1, \dots, f_T) = \sum_i \|y_i - \hat{y}_i\|^2 + \sum_t \Omega(f_t)$$

TREE BOOSTING IN A NUTSHELL

Model is trained in a sequential / additive manner:

$$\mathcal{L}^{(t)} = \sum_{i=1}^n \ell(y_i, \hat{y}_i^{(t-1)} + f_t(x_i)) + \Omega(f_t)$$

Recap:

$$v(y + \epsilon) \approx v(y) + \epsilon v'(y) + \frac{\epsilon^2}{2} v''(y)$$

Taylor approximation (order 2):

$$\mathcal{L}^{(t)} \approx \sum_{i=1}^n \left[\ell(y_i, \hat{y}_i^{(t-1)}) + f_t(x_i) g_i + \frac{f_t^2(x_i)}{2} h_i \right] + \Omega(f_t)$$

$$g_i = \partial_{\hat{y}^{(t-1)}} \ell(y_i, \hat{y}^{(t-1)}) \quad (\text{gradient})$$

$$h_i = \partial_{\hat{y}^{(t-1)}}^2 \ell(y_i, \hat{y}^{(t-1)}) \quad (\text{hessian})$$

TREE BOOSTING IN A NUTSHELL

$$\mathcal{L}^{(t)} \approx \sum_{i=1}^n \left[\ell(y_i, \hat{y}_i^{(t-1)}) + f_t(x_i)g_i + \frac{f_t^2(x_i)}{2}h_i \right] + \Omega(f_t)$$

Example with MSE / L2 loss:

$$\ell(y_i, \hat{y}_i^{(t-1)}) = \|y_i - \hat{y}_i^{(t-1)}\|^2$$

$$g_i = \partial_{\hat{y}_i^{(t-1)}} \ell(y_i, \hat{y}_i^{(t-1)}) = -2(y_i - \hat{y}_i^{(t-1)})$$

$$h_i = \partial_{\hat{y}_i^{(t-1)}}^2 \ell(y_i, \hat{y}_i^{(t-1)}) = 2$$

TREE BOOSTING IN A NUTSHELL

$$\mathcal{L}^{(t)} \approx \sum_{i=1}^n \left[\ell(y_i, \hat{y}_i^{(t-1)}) + f_t(x_i)g_i + \frac{f_t^2(x_i)}{2}h_i \right] + \Omega(f_t)$$

Removing constant terms we just need to minimize w.r.t. f_t :

$$\tilde{\mathcal{L}}^{(t)} = \sum_{i=1}^n \left[f_t(x_i)g_i + \frac{f_t^2(x_i)}{2}h_i \right] + \Omega(f_t)$$

TREE BOOSTING IN A NUTSHELL

$$\tilde{\mathcal{L}}^{(t)} = \sum_{i=1}^n \left[f_t(x_i)g_i + \frac{f_t^2(x_i)}{2}h_i \right] + \Omega(f_t)$$

which can be rewritten:

$$\begin{aligned} \tilde{\mathcal{L}}^{(t)} &= \sum_{i=1}^n \left[f_t(x_i)g_i + \frac{f_t^2(x_i)}{2}h_i \right] + \gamma K + \frac{\lambda}{2} \sum_{k=1}^K (z_k^{f_t})^2 \\ &= \sum_{k=1}^K \left[\left(\sum_{i \in I_k} g_i \right) z_k^{f_t} + \frac{1}{2} \left(\lambda + \sum_{i \in I_k} h_i \right) (z_k^{f_t})^2 \right] + \gamma K \end{aligned}$$

where: $I_k = \{i \mid q(x_i) = k\}$ (samples in leaf j)

TREE BOOSTING IN A NUTSHELL

Minimizing this:

$$\tilde{\mathcal{L}}^{(t)} = \sum_{k=1}^K \left[\left(\sum_{i \in I_k} g_i \right) z_k^{f_t} + \frac{1}{2} \left(\lambda + \sum_{i \in I_k} h_i \right) (z_k^{f_t})^2 \right] + \gamma K$$

leads to:
$$z_k^* = - \frac{\sum_{i \in I_k} g_i}{\lambda + \sum_{i \in I_k} h_i}$$

and:
$$\tilde{\mathcal{L}}^{(t)} = - \frac{1}{2} \sum_{k=1}^K \frac{(\sum_{i \in I_k} g_i)^2}{\lambda + \sum_{i \in I_k} h_i} + \gamma K$$

TREE BOOSTING IN A NUTSHELL

Greedy optimization of:

$$\tilde{\mathcal{L}}^{(t)} = -\frac{1}{2} \sum_{k=1}^K \frac{(\sum_{i \in I_k} g_i)^2}{\lambda + \sum_{i \in I_k} h_i^2} + \gamma K$$

For one leaf splitting it leads to the following criteria that should be maximized:

$$\mathcal{C}_{\text{split}} = \frac{1}{2} \left[\frac{(\sum_{i \in I_L} g_i)^2}{\lambda + \sum_{i \in I_L} h_i^2} + \frac{(\sum_{i \in I_R} g_i)^2}{\lambda + \sum_{i \in I_R} h_i^2} - \frac{(\sum_{i \in I} g_i)^2}{\lambda + \sum_{i \in I} h_i^2} \right] - \gamma$$

It corresponds to the loss reduction by splitting: $I = I_L \cup I_R$

TREE BOOSTING IN A NUTSHELL

Criteria that should be maximized:

$$\mathcal{C}_{\text{split}} = \frac{1}{2} \left[\frac{(\sum_{i \in I_L} g_i)^2}{\lambda + \sum_{i \in I_L} h_i^2} + \frac{(\sum_{i \in I_R} g_i)^2}{\lambda + \sum_{i \in I_R} h_i^2} - \frac{(\sum_{i \in I} g_i)^2}{\lambda + \sum_{i \in I} h_i^2} \right] - \gamma$$

Example with MSE / L2 loss: $\ell_i(y_i, \hat{y}_i) = \|y_i - \hat{y}_i\|^2$

$$\mathcal{C}_{\text{split}} = \frac{1}{2} \left[\frac{(\sum_{i \in I_L} 2(y_i - \hat{y}_i^{(t-1)}))^2}{\lambda + \sum_{i \in I_L} 4} + \frac{(\sum_{i \in I_R} 2(y_i - \hat{y}_i^{(t-1)}))^2}{\lambda + \sum_{i \in I_R} 4} - \frac{(\sum_{i \in I} 2(y_i - \hat{y}_i^{(t-1)}))^2}{\lambda + \sum_{i \in I} 4} \right] - \gamma$$

Remark: It's close to a variance impurity criterion but not equal

TREE BOOSTING ALGORITHM

Algorithm 1: Exact Greedy Algorithm for Split Finding

Input: I , instance set of current node

Input: d , feature dimension

$gain \leftarrow 0$

$G \leftarrow \sum_{i \in I} g_i, H \leftarrow \sum_{i \in I} h_i$

for $k = 1$ **to** m **do**

$G_L \leftarrow 0, H_L \leftarrow 0$

for j **in** $sorted(I, \text{by } \mathbf{x}_{jk})$ **do**

$G_L \leftarrow G_L + g_j, H_L \leftarrow H_L + h_j$

$G_R \leftarrow G - G_L, H_R \leftarrow H - H_L$

$score \leftarrow \max(score, \frac{G_L^2}{H_L + \lambda} + \frac{G_R^2}{H_R + \lambda} - \frac{G^2}{H + \lambda})$

end

end

Output: Split with max score

GOING BEYOND:

- Approximate splitting (feature binning)
- Parallel implementation (multi-thread & multi-machine)
- Sparsity aware split finding (think one-hot encoding)
- Cache-aware access
- Monotonic constraints

APPROXIMATE SPLITTING

Algorithm 2: Approximate Algorithm for Split Finding

```
for  $k = 1$  to  $m$  do  
    | Propose  $S_k = \{s_{k1}, s_{k2}, \dots, s_{kl}\}$  by percentiles on feature  $k$ .  
    | Proposal can be done per tree (global), or per split(local).  
end  
for  $k = 1$  to  $m$  do  
    |  $G_{kv} \leftarrow \sum_{j \in \{j | s_{k,v} \geq \mathbf{x}_{jk} > s_{k,v-1}\}} g_j$   
    |  $H_{kv} \leftarrow \sum_{j \in \{j | s_{k,v} \geq \mathbf{x}_{jk} > s_{k,v-1}\}} h_j$   
end
```

Follow same step as in previous section to find max score only among proposed splits.

Questions?

GBRT HANDS ON



See: code in `tinygbt.py` folder